

Aadith Lasar

LY CORE-2

2203262

MACHINE LEARNING LAB-2

Title: Download the any dataset from UCI or Data.org or from any other data repositories and perform the basic data pre-processing steps using python/R .

After completion of this experiment students will be able to:

- - Learn to pre-process dataset
- Learn to use pandas and sklearn

Aim: To download the any dataset from UCI or Data.org or from any other data repositories and perform the basic data pre-processing steps using python/R .

Theory:

Step 1 : Import the libraries

Step 2 : Import the data-set

Step 3 : Check out the missing values

Step 4 : See the Categorical Values

Step 5 : Splitting the data-set into Training and Test Set

Data cleaning:

The main aim of Data Cleaning is to identify and remove errors & duplicate data, in order to create a reliable dataset. This improves the quality of the training data for analytics and enables accurate decision-making.

Needless to say, data cleansing is a time-consuming process and most data scientists spend an enormous amount of time in enhancing the quality of the data. However, there are various methods to identify and classify data for data cleansing

There are mainly two distinct techniques, namely Qualitative and Quantitative techniques to classify data errors. Qualitative techniques involve rules, constraints, and patterns to identify errors.

On the other hand, Quantitative techniques employ statistical techniques to identify errors in the trained data.

Normalisation:

Normalization is a scaling technique in which values are shifted and rescaled so that they end up ranging between 0 and 1. It is also known as Min-Max scaling.

Here's the formula for normalization:

$$X' = \frac{X - X_{min}}{X_{max} - X_{min}}$$

Here, Xmax and Xmin are the maximum and the minimum values of the feature respectively.

- When the value of X is the minimum value in the column, the numerator will be 0, and hence X' is 0
- On the other hand, when the value of X is the maximum value in the column, the numerator is equal to the denominator and thus the value of X' is 1
- If the value of X is between the minimum and the maximum value, then the value of X' is between 0 and 1

Standardisation:

Standardization is another scaling technique where the values are centered around the mean with a unit standard deviation. This means that the mean of the attribute becomes zero and the resultant distribution has a unit standard deviation.

Here's the formula for standardization:

$$X' = \frac{X - \mu}{\sigma}$$

μ is the mean of the feature values and σ is the standard deviation of the feature values. Note that in this case, the values are not restricted to a particular range.

`X_train,X_test,y_train,y_test= train_test_split(X,y,test_size=0.2,random_state= 0)`

Split arrays or matrices into random train and test subsets Quick utility that wraps input validation and `next(ShuffleSplit().split(X, y))` and application to input data into a single call for splitting (and optionally subsampling) data in a oneliner.

`Imputer(missing_values='NaN', strategy='mean', axis=0)`

Imputation transformer for completing missing values.

`pandas.read_csv()`

Read a comma-separated values (csv) file into DataFrame.

Also supports optionally iterating or breaking of the file into chunks.

Program:-

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

data = pd.read_csv("Whole_cust.csv")

data.shape

data.isnull().sum()

data.dropna()

data.skew()

data.kurtosis()

import seaborn as sns

sns.boxplot(x="Channel", y="Region", data=data)

data.plot(kind="scatter", x="Grocery", y="Detergents_Paper")
```

```

sns.violinplot(x="Channel", y="Region", data=data)\

sns.pairplot(data, hue="Fresh", height=4)

data.corr()

from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.model_selection import train_test_split

X = data.iloc[:, :-1].values
y = data.iloc[:, -1].values
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 0)

from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

from sklearn.naive_bayes import GaussianNB
nvclassifier = GaussianNB()
nvclassifier.fit(X_train, y_train)

y_pred = nvclassifier.predict(X_test)
print(y_pred)

y_compare = py.vstack((y_test,y_pred)).T
y_compare[:5,:]

from sklearn.metrics import accuracy_score, confusion_matrix

acc = accuracy_score(y_test, y_pred) * 100
print("Accuracy =", acc)

confusion_matrix(y_test, y_pred)

from sklearn.metrics import classification_report

print(classification_report(y_test,y_pred))

```

Output:-

```
jupyter Wholesale_customers Last Checkpoint: 25 minutes ago (unsaved changes)
File Edit View Insert Cell Kernel Help Trusted Python 3 (pykernel)
In [1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
#AADITH LASAR
#2203262 LY-CORE-2
In [2]: data = pd.read_csv("Whole_cust.csv")
In [3]: data
Out[3]:
```

	Channel	Region	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicassen
0	2	3	12669	9656	7561	214	2674	1338
1	2	3	7057	9810	9568	1762	3293	1776
2	2	3	6353	8808	7684	2405	3516	7844
3	1	3	13265	1196	4221	6404	507	1788
4	2	3	22615	5410	7198	3915	1777	5185
...
435	1	3	29703	12051	16027	13135	182	2204
436	1	3	39228	1431	764	4510	93	2346
437	2	3	14531	15488	30243	437	14841	1867
438	1	3	10290	1981	2232	1038	168	2125
439	1	3	2787	1698	2510	65	477	52

440 rows x 8 columns

```
In [4]: data.shape
Out[4]: (440, 8)
In [5]: data.isnull().sum()
Out[5]: Channel      0
Region            0
Fresh             0
Milk              0
Grocery           0
Frozen            0
Detergents_Paper  0
Delicassen        0
dtype: int64
In [6]: data.dropna()
Out[6]:
```

	Channel	Region	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicassen
0	2	3	12669	9656	7561	214	2674	1338
1	2	3	7057	9810	9568	1762	3293	1776
2	2	3	6353	8808	7684	2405	3516	7844
3	1	3	13265	1196	4221	6404	507	1788
4	2	3	22615	5410	7198	3915	1777	5185
...
435	1	3	29703	12051	16027	13135	182	2204
436	1	3	39228	1431	764	4510	93	2346
437	2	3	14531	15488	30243	437	14841	1867
438	1	3	10290	1981	2232	1038	168	2125
439	1	3	2787	1698	2510	65	477	52

440 rows x 8 columns


```
In [7]: data.skew()
Out[7]: Channel      0.760951
Region    -1.283627
Fresh       2.561323
Milk        4.053755
Grocery     3.587429
Frozen      5.907986
Detergents_Paper  3.631851
Delicassen  11.151586
```

Run Code

```
In [9]: import seaborn as sns
```


```
In [18]: sns.boxplot(x="Channel", y="Region", data=data)
```

```
Out[18]: <Axes: xlabel='Channel', ylabel='Region'>
```




```
In [19]: data.plot(kind="scatter", x="Grocery", y="Detergents_Paper")
```

```
Out[19]: <Axes: xlabel='Grocery', ylabel='Detergents_Paper'>
```



```
In [20]: sns.violinplot(x="Channel", y="Region", data=data)
```

```
Out[20]: <Axes: xlabel='Channel', ylabel='Region'>
```



Run Code

```
In [22]: sns.pairplot(data, hue="Fresh", height=4)
```

```
Out[22]: <seaborn.axisgrid.PairGrid at 0x2a6622490>
```



```
In [23]: data.corr()
```

```
Out[23]:
```

	Channel	Region	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicassen
Channel	1.000000	0.062028	-0.169172	0.460720	0.608792	-0.202046	0.636026	0.056011
Region	0.062028	1.000000	0.055287	0.032288	0.007696	-0.021044	-0.001483	0.045212
Fresh	-0.169172	0.055287	1.000000	0.100510	-0.011854	0.345881	-0.101953	0.244690
Milk	0.460720	0.032288	0.100510	1.000000	0.728335	0.123994	0.661816	0.406368
Grocery	0.608792	0.007696	-0.011854	0.728335	1.000000	-0.040193	0.924641	0.205497
Frozen	-0.202046	-0.021044	0.345881	0.123994	-0.040193	1.000000	-0.131525	0.390947
Detergents_Paper	0.636026	-0.001483	-0.101953	0.661816	0.924641	-0.131525	1.000000	0.069291
Delicassen	0.056011	0.045212	0.244690	0.406368	0.205497	0.390947	0.069291	1.000000

```
In [24]: from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.model_selection import train_test_split
```

```
In [25]: X = data.iloc[:, :-1].values
y = data.iloc[:, -1].values
```

Jupyter Wholesale_customers Last Checkpoint: 30 minutes ago (autosaved) Logout

File Edit View Insert Cell Kernel Help Trusted Python 3 (ipykernel)

```

In [26]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

In [27]: from sklearn.naive_bayes import GaussianNB
nvclassifier = GaussianNB()
nvclassifier.fit(X_train, y_train)

Out[27]: GaussianNB
GaussianNB()

In [28]: y_pred = nvclassifier.predict(X_test)
print(y_pred)

[2563  46  548  834 405  436 156 1215 405  548  46 2563 395  548
  548 405  46  3 1117 395 1215  3 2563 548 750 405 2563 548
  3 2563 395  46  3 395 548  46  46 247 395 548 395 548
  548  46 548  46  46 1117 405 548 395 548 2563 436  46 1117
 379 548 436 548 548 436 548  46  3 120  3 436 750  3
 405 1117  46 405 548 548 405 379 395 548  3 548 834 395
  46 436  46 610]

In [29]: y_compare = py.vstack((y_test,y_pred)).T
y_compare[:5,:]
```

```

Out[29]: array([[ 806, 2563],
 [ 445,  46],
 [ 291,  548],
 [1625,  834],
 [1333, 405]])

In [30]: from sklearn.metrics import accuracy_score, confusion_matrix

acc = accuracy_score(y_test, y_pred) * 100
print("Accuracy =", acc)

Accuracy = 0.0

In [31]: confusion_matrix(y_test, y_pred)

Out[31]: array([[0, 0, 0, ..., 0, 0, 0],
 [0, 0, 1, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 ...,
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 [1, 0, 0, ..., 0, 0, 0]])

In [32]: from sklearn.metrics import classification_report

print(classification_report(y_test,y_pred))
```

```

              precision    recall  f1-score   support

 2563          0.00         0.00         0.00         1.0
 2602          0.00         0.00         0.00         1.0
 2784          0.00         0.00         0.00         1.0
 2931          0.00         0.00         0.00         1.0
 3029          0.00         0.00         0.00         1.0
 5137          0.00         0.00         0.00         1.0
 5609          0.00         0.00         0.00         1.0
 6250          0.00         0.00         0.00         1.0

 accuracy          0.00         0.00         0.00         88.0
 macro avg          0.00         0.00         0.00         88.0
 weighted avg          0.00         0.00         0.00         88.0

/Users/aadithlasar/miniforge3/lib/python3.9/site-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  warn_prf(average, modifier, msg_start, len(result))
/Users/aadithlasar/miniforge3/lib/python3.9/site-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division`

```

In []:

Conclusion:

Thus we have successfully implemented pre-processing operations on a dataset